

Pilot results

04 OCT 2026 | VERSION 1.0

Study setup

This fictional Harbour AI pilot used 200 synthetic support tickets from five simulated shops. Tickets covered delivery questions, order lookups, and refund-policy questions. Two reviewers checked each draft against the approved reference material. These invented results are for an AskBase demonstration, not a real product benchmark.

Targets versus actual results

For Harbour AI, the handling-time reduction target was at least 40%. The measured median fell from 10 minutes to 7 minutes: a 30% reduction, missing the target by 10 percentage points. Citation correctness was 96%, above the 95% target. Unsupported answers were 3%, above the required below-2% threshold. Of the three release targets, only citation correctness passed.

Failure patterns

Four of the six unsupported Harbour AI drafts relied on outdated refund-policy text. The other two mixed up similarly named products. Reviewers also found that status questions became difficult when the order lookup tool timed out. Tool timeouts must produce an explicit failure message rather than an invented order status.

Release decision

Harbour AI is approved only for a human-reviewed beta with 10 shops. Unattended sending remains blocked because handling-time reduction and unsupported-answer rates missed their targets. Human review remains mandatory even though citation correctness passed. This report does not approve Arabic support, voice support, autonomous refunds, or automatic order changes.

Next experiment

Before reassessing unattended sending, Harbour AI will add document-expiry checks, clearer product identifiers, and explicit tool-timeout handling. The next evaluation must contain at least 100 new refund-policy tickets, separate from the first pilot set. The owner must recheck all three release targets; improved citations alone do not clear the release gate.